



Introduction to Database Systems: CS312

MongoDB

Oliver BONHAM-CARTER

26 Oct 2020

The Problem with SQL

The Problem With SQL

About Mongo

Schema

Language

Start MongoDB

Killing a container

Collections

Finding Information

A Simple Base

Queries

Shutting Down

Queries

Practical 03

Consider this

firstName	lastName	primaryAddr

- Let's say that we have a (perfectly) working SQL table
- The schema has been designed and coded for *current data requirements*

Table Update (i)

The Problem
With SQL

About Mongo

Schema

Language

Start
MongoDB

Killing a
container

Collections

Finding
Information

A Simple
Base

Queries

Shutting
Down

Queries

Practical 03

Consider this

firstName	lastName	primaryAddr	secondAddr

- The data we collect has changed.
- We need to update our schema for the *new data requirements*

Table Update (ii)

The Problem
With SQL

About Mongo

Schema

Language

Start
MongoDB

Killing a
container

Collections

Finding
Information

A Simple
Base

Queries

Shutting
Down




















Queries

Practical 03

Consider this

firstName	lastName	primaryAddr	secondAddr	thirdAddr

- Our needs have changed again and the SQL table must be updated.
- The schema is reprogrammed

firstName	lastName	primaryAddr	secondAddr	thirdAddr
				
				
				
				

- We expect that the table will be full when in use
- Expectations are not always fulfilled...



Consider this

- But, in reality, much of the table is empty!
- The table can easily get huge and be hard to manage.

We Might Stop and Ask Ourselves...

The Problem
With SQL

About Mongo

Schema

Language

Start
MongoDB

Killing a
container

Collections

Finding
Information

A Simple
Base


Queries

Shutting
Down

Queries

Practical 03

Consider this



- What can we do to stop having to redesign our database schema with our changing data?
- Is SQL the right type of database management system for our changing data requirements?

The Problem With SQL

About Mongo

Schema

Language

Start MongoDB

Killing a container

Collections

Finding Information

A Simple Base

Queries

Shutting Down

Queries

Practical 03

Consider this



mongoDB®

- NoSQL: *Not Only SQL* database systems that support SQL-like query languages, but are used increasingly in big data applications and real-time web applications.
- The stored data is allowed to change

● <https://www.mongodb.com/>

Database Language Guide

SQL systems versus NoSQL

The Problem
With SQL

About Mongo

Schema

Language

Start
MongoDB

Killing a
container

Collections

Finding
Information

A Simple
Base

Queries

Shutting
Down

Queries

Practical 03

Consider this

RDBMS		MongoDB
Database	⇒	Database
Table, View	⇒	Collection
Row	⇒	Document (BSON)
Column	⇒	Field
Index	⇒	Index
Join	⇒	Embedded Document
Foreign Key	⇒	Reference
Partition	⇒	Shard

- The terms are different but their meanings are similar
- *Schema-less*, collections (like tables) are populated by any data
- *Documents* are similar to the *tuples* of Sqlite3 programming



Schema Free

The Problem
With SQL

About Mongo

Schema

Language

Start
MongoDB

Killing a
container

Collections

Finding
Information

A Simple
Base


Queries

Shutting
Down

Queries

Practical 03

Consider this



Collection

- No pre-defined data schema
 - Data may be entered in absence of a defined schema
- **Documents** (*rows*) of **collections** (*DB's*) may have different types of data

Schema Free

Mostly similar documents

The Problem
With SQL

About Mongo

Schema

Language

Start
MongoDB

Killing a
container

Collections

Finding
Information

A Simple
Base

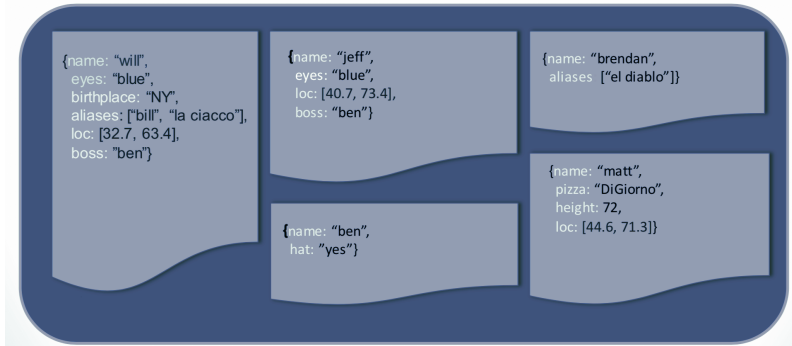
Queries

Shutting
Down

Queries

Practical 03

Consider this



- Sometimes not all the data is available to create a document.
- The query *interprets* missing data as NULL entries



JSON and MongoDB Code

The Problem
With SQL

About Mongo

Schema

Language

JSON

Start
MongoDB

Killing a
container

Collections

Finding
Information

A Simple
Base

Queries

Shutting
Down

Queries

Practical 03

Consider this

- **Data is in name / value pairs**

- A name/value pair consists of a field name followed by a colon, followed by a value:

- Example: `{ "name": "R2-D2" }`

- **Data is separated by commas**

- Example: `{ "name": "R2-D2", race : "Droid" }`

- **Curly braces hold objects**

- Example: `{ "name": "R2-D2", race : "Droid", affiliation: "rebels" }`

- **An array is stored in brackets []**

- Example `[{ "name": "R2-D2", race : "Droid", affiliation: "rebels" }, { "name": "Yoda", affiliation: "rebels" }]`

The Problem
With SQL

About Mongo

Schema

Language

JSON

Start
MongoDB

Killing a
container

Collections

Finding
Information

A Simple
Base

Queries

Shutting
Down

Queries

Practical 03

Consider this

We will see more of these commands later!

- **Db.collection** specifies the collection or the *table* in which to store the document (*tuple*)
- **Create**
 - `db.collection.insert()`
 - `db.collection.save()`
 - `db.collection.update()`
- **Read**
 - `db.collection.find()`
 - `db.collection.findOne()`
- **Update**
 - `db.collection.update()`
- **Delete**
 - `db.collection.remove()`



Getting started with Mongo in Docker

The Problem
With SQL

About Mongo

Schema

Language

Start
MongoDB

Windows
MacOS and
Linux

Killing a
container

Collections

Finding
Information

A Simple
Base

Queries

Shutting
Down

Queries

Practical 03

Consider this
13 / 33

Create a Docker container from DockerHub for Mongo: Windows.

```
docker pull mongo
```

Create a Docker container from DockerHub for Mongo: MacOS and Linux.

```
sudo docker pull mongo
```

Setup Mongo in Docker Container

Windows, Powershell

The Problem
With SQL

About Mongo

Schema

Language

Start
MongoDB

Windows
MacOS and
Linux

Killing a
container

Collections

Finding
Information

A Simple
Base

Queries

Shutting
Down

Queries

Practical 03

Consider this
14 / 33



Create a local directory for your data

```
mkdir c:\mongodata
```

Start a Docker container running the Mongo DB server

```
docker run -it -v c:\mongodata:/data/db --name mongodb -d mongo
```

Check log to see that the server is operational

```
docker logs mongodb  
docker ps
```

Run instance of MongoDB, goes into root of container.

```
docker exec -it mongodb bash
```

Setup Mongo in Docker Container

Windows, Powershell

The Problem
With SQL

About Mongo

Schema

Language

Start
MongoDB

Windows
MacOS and
Linux

Killing a
container

Collections

Finding
Information

A Simple
Base

Queries

Shutting
Down

Queries

Practical 03

Consider this
15 / 33



Start the MongoDB client

```
mongo
```

You are now able to run MongoDB commands here.

Note: exit to quit.

Setup Mongo in Docker Container

Windows, Powershell

The Problem
With SQL

About Mongo

Schema

Language

Start
MongoDB

Windows

MacOS and
Linux

Killing a
container

Collections

Finding
Information

A Simple
Base

Queries

Shutting
Down

Queries

Practical 03

Consider this
16 / 33



Leave the container

```
exit
```

Stop Mongo container

```
docker stop mongodb
```

Removing all stopped containers, if necessary due to errors in
launching container

```
docker rm $(docker ps -a -q)
```

Setup Mongo in Docker Container

MacOS and Linux, Terminal

The Problem
With SQL

About Mongo

Schema

Language

Start
MongoDB

Windows
MacOS and
Linux

Killing a
container

Collections

Finding
Information

A Simple
Base

Queries

Shutting
Down

Queries

Practical 03

Consider this
17 / 33



Create a directory for data to persist.

```
mkdir -p ~/mongodata
```

Start a Docker container running the Mongo DB server

```
sudo docker run -it -v ~/mongodata:/data/db --name mongodb -d mongo
```

Check log to see that the server is operational

```
sudo docker logs mongodb  
sudo docker ps
```

Run instance of MongoDB, goes into root of container.

```
sudo docker exec -it mongodb bash
```

Setup Mongo in Docker Container

MacOS and Linux, Terminal

The Problem
With SQL

About Mongo

Schema

Language

Start
MongoDB

Windows

MacOS and
Linux

Killing a
container

Collections

Finding
Information

A Simple
Base

Queries

Shutting
Down

Queries

Practical 03

Consider this
18 / 33



Start the MongoDB client

```
mongo
```

You are now able to run MongoDB commands here.

Note: exit to quit.

Setup Mongo in Docker Container

MacOS and Linux, Terminal

The Problem
With SQL

About Mongo

Schema

Language

Start
MongoDB

Windows
MacOS and
Linux

Killing a
container

Collections

Finding
Information

A Simple
Base

Queries

Shutting
Down

Queries

Practical 03

Consider this
19 / 33



Leave the container

```
exit
```

Stop MongoDB container

```
sudo docker stop mongodb
```

Removing all stopped containers, if necessary due to errors in
launching container

```
sudo docker rm $(docker ps -a -q)
```

Killing a Docker Container

If you have trouble running it, try relaunching the container

Find out the container's ID

```
docker ps
```

Your container's ID. This one is; 4e9d863407ac

CONTAINER ID
4e9d863407ac

IMAGE
mongo

Stop MongoDB container

```
docker stop mongodb # Windows
```

```
sudo docker stop mongodb # MacOS and Linux
```

Remove MongoDB container

```
docker rm 4e9d863407ac # Windows
```

```
sudo docker rm 4e9d863407ac # MacOS and Linux
```

You should now be able to run your run and execute commands from above.



The Problem
With SQL

About Mongo

Schema

Language

Start
MongoDB

Killing a
container

Collections

Finding
Information

A Simple
Base

Queries

Shutting
Down

Queries

Practical 03

Consider this

Simple Insert into the *Greetings* collection

```
db.Greetings.drop()
```

```
db.Greetings.insert({"Hello":"English greeting"})
```

```
db.Greetings.insert({"Goodbye":"English farewell"})
```

```
db.Greetings.insert({"Bonjour":"French greeting"})
```

```
db.Greetings.insert({"Au revoir":"French farewell"})
```

```
db.Greetings.insert({"Buongiorno":"Italian greeting"})
```

```
db.Greetings.insert({"Arrivederci":"Italian farewell"})
```

Find contents of *Greetings* collection:

```
select * from Greetings;
```

```
db.Greetings.find({}, {})
```

```
db.Greetings.find({}, {}).pretty();
```

```
db.Greetings.find({}, {_id:0}).pretty();
```



General Finds in a "Music" Collection

For example ... true or 1

The Problem
With SQL

About Mongo

Schema

Language

Start
MongoDB

Killing a
container

Collections

Finding
Information

A Simple
Base

Queries

Shutting
Down

Queries

Practical 03

Consider this

- The `find()` method with no parameters returns all documents from a collection and returns all fields for the documents.

The following operation returns all documents in the `music` collection

```
db.music.find({}, {})
```

Returns all "bandName" in the `music` collection

```
db.music.find({ }, {"bandName":True})  
db.music.find({ }, {"bandName":1})
```



General Finds in a “Music” Collection

For example ... true or 1

The Problem
With SQL

About Mongo

Schema

Language

Start
MongoDB

Killing a
container

Collections

Finding
Information

A Simple
Base

Queries

Shutting
Down

Queries

Practical 03

Consider this

```
db.collection.find({ Search field }, { What to filter in field })
```

- The `find()` method can also be used to specify particular fields to show from a collection-wide search.

The following operation returns all records with no filtration

```
db.music.find({}, {})
```

The following operation filters-out only the `bandName` from the `music` collection

```
db.music.find({ }, {"bandName":true})
```

Returns all `bandName` and songs in the `music` collection

```
db.music.find({ }, {"bandName":true, "songs":true})
```


Insert into the *Country* collection



```
db.Country.drop()
db.Country.insert({
  language_en: "English",
  language_fr: "anglais",
  country_en: "England",
  country_fr : "l'angleterre"
})
```

Adding to the Collection

The Problem
With SQL

About Mongo

Schema

Language

Start
MongoDB

Killing a
container

Collections

Finding
Information

A Simple
Base

Queries

Shutting
Down

Queries

Practical 03

Consider this

```
db.Country.insert({  
  language_en: "French",  
  language_fr: "français",  
  country_en: "France",  
  country_fr : "la France"  
})
```

```
db.Country.insert({  
  language_en: "Italian",  
  language_fr: "italien",  
  country_en: "Italy",  
  country_fr : "l'Italie"  
})
```

Simple Queries

The Problem
With SQL

About Mongo

Schema

Language

Start
MongoDB

Killing a
container

Collections

Finding
Information

A Simple
Base

Queries

Shutting
Down

Queries

Practical 03

Consider this

Find contents of *Country* collection:
`select * from Country;`

```
db.Country.find()  
db.Country.find().pretty()
```

```
> db.Country.find().pretty()  
{  
  "_id" : ObjectId("5c89c87211455eb94061b2f1"),  
  "language_en" : "English",  
  "language_fr" : "anglais",  
  "country_en" : "England",  
  "country_fr" : "l'angleterre"  
}  
{  
  "_id" : ObjectId("5c89c87211455eb94061b2f2"),  
  "language_en" : "French",  
  "language_fr" : "français",  
  "country_en" : "France",  
  "country_fr" : "la France"  
}  
{  
  "_id" : ObjectId("5c89c87311455eb94061b2f3"),  
  "language_en" : "Italian",  
  "language_fr" : "italien",  
  "country_en" : "Italy",  
  "country_fr" : "l'Italie"  
}
```



Simple Queries

The Problem
With SQL

About Mongo

Schema

Language

Start
MongoDB

Killing a
container

Collections

Finding
Information

A Simple
Base

Queries

Shutting
Down

Queries

Practical 03

Consider this

Find specific contents of *Country* collection

```
select * from Country where country_en == "France";
```

```
db.Country.find({country_en : "France"}).pretty()
```

Find specific contents of *Country* collection

```
select * from Country where language_en == "English";
```

```
db.Country.find({language_en : "English"}).pretty()
```

Find specific contents of *Country* collection

```
select language_fr, language_en from Country where language_en ==  
"English";
```

```
db.Country.find({language_en : "English"},  
  {_id:0}).pretty()
```

```
db.Country.find({language_en:"English",  
  {language_fr : 1, language_en : 1}).pretty()
```



How to shut down a session

The Problem
With SQL

About Mongo

Schema

Language

Start
MongoDB

Killing a
container

Collections

Finding
Information

A Simple
Base

Queries

Shutting
Down

Queries

Practical 03

Consider this

Drop the collection (Greetings): Destroy the data, remove collection

```
db.Greetings.drop()
```

Drop the collection (Country): Destroy the data, remove collection

```
db.Country.drop()
```

Closing down

- Type `exit` in the client terminal

Advanced Data and Queries

Let's horse around with data and code ...



Locate your data file, open it in Atom and copy and past the contents (all at once) into the MondoDB client

```
syntheticData_small.json
```

Synthetic data available from:

<https://next.json-generator.com/>

The Problem
With SQL

About Mongo

Schema

Language

Start
MongoDB

Killing a
container

Collections

Finding
Information

A Simple
Base

Queries

Shutting
Down

Queries

Practical 03

Consider this

The Problem
With SQL

About Mongo

Schema

Language

Start
MongoDB

Killing a
container

Collections

Finding
Information

A Simple
Base

Queries

Shutting
Down

Queries

Practical 03

Consider this



```
db.collection.find({searchSpace},{showAttrib:1})
```

What do the following queries do?

```
db.employee.find({},{"name.last":1, "name.first":1})
db.employee.find({},{company:1})
db.employee.find({},{"name.last":1,address:1})
db.employee.find({},{company:1,registered:1,_id:0})
db.employee.find({},{company:1,"friends.name":1,_id:0})
db.employee.find({company:"RADIANTIX"},
{"friends.name":1,_id:0})
```

The Problem
With SQL

About Mongo

Schema

Language

Start
MongoDB

Killing a
container

Collections

Finding
Information

A Simple
Base

Queries

Shutting
Down

Queries

Practical 03

Consider this

```
db.collection.find({searchSpace},{showAttrib:1})
```

Rewrite the queries to get...

- The last name, first name and balance information of each document
- The company name, latitude and longitude of each document
- The list of friends of record associated with the company PROSURE
- The phone number for the PROSURE company



GitHub submission

- Write queries in MongoDB
- File: `sandbox/03_practical/queries.md`
 - Move the directory (`03_practical/`) into your practicals repository (with files) to submit your work.
- GitHub Classroom repository link:
<https://classroom.github.com/a/3FKw8tMD>
- Useful commands:
 - `git add -A`
 - `git commit -m "Your commit caption here"`
 - `git push`
- Your check-marked assignment is due by 2nd Nov 2020, 11:30 EDT

THINK



Consider this...

The Problem
With SQL

About Mongo

Schema

Language

Start
MongoDB

Killing a
container

Collections

Finding
Information

A Simple
Base

Queries

Shutting
Down

Queries

Practical 03

Consider this



THINK

- Can you create and populate a new Mongo database?
- Can you write sophisticated queries in your database?